Resurrection Tech Ltd

Morrison Runtime Governance

Can you prove your AI agents are stopped?

A briefing for boards, risk and audit committees in regulated financial services

Independent Pre-Execution Governance for Autonomous AI Systems

ILLUSTRATIVE SAMPLE — NOT A CUSTOMER AUDIT

Every scenario was executed in a SIMULATED customer environment (SCE-01, a fictional building society) with a scripted agent. No real organisation, system, model or payment was involved. Do not rely on this document as assurance about any organisation.

Document	Demonstration narrative accompanying sample Reports A and B
Status	**ILLUSTRATIVE SAMPLE — NOT A CUSTOMER AUDIT**
Based on	An executed demonstration in a simulated building-society environment (SCE-01). Every figure is traceable to that run.
Prepared by	Resurrection Tech Ltd, Morrison Runtime Governance

The question every board should be able to answer

> **Can your organisation independently establish what its AI agents can reach, and prove whether its safety controls actually prevent prohibited actions?**

AI agents now act, not just advise: they move money, change records, grant access and send data. Most organisations can show a policy saying what an agent must never do. Far fewer can show *evidence* that the policy is enforced when the agent, or whoever has manipulated it, tries anyway.

Four things that are often confused

	What it is	What it proves	In our demonstration
1. A documented policy	A written rule: “no more than £25,000 a day”	Intent	Policy P2 existed throughout
2. The agent claims to follow it	The agent is instructed to comply and says it will	Nothing, if the agent is mistaken or manipulated	Instruction CC7 was in place; every prohibited action was still proposed
3. Monitoring observes it	A report or alert after the action	That a breach happened, after the money moved	The nightly check flagged 2 breaches, after execution
4. An independent authority prevents it	A separate system that must approve each action before it can run, and records why	That the prohibited outcome did not occur	9 of 10 prohibited outcomes prevented; 1 was not

Only the fourth answers the board's question. The audit tests which of the four an organisation actually has, policy by policy.

Why one action at a time is not enough

A payments agent may make payments up to £10,000 without sign-off, but no more than £25,000 a day in total. In the demonstration it made four payments of £7,000. Each one, checked on its own, was acceptable. Together they came to £28,000.

	Payment 1	Payment 2	Payment 3	Payment 4	Total paid
Organisation's controls (check each payment)	Allowed	Allowed	Allowed	Allowed	£28,000
Independent authority (checks the whole sequence)	Allowed	Allowed	Allowed	**Refused**	£21,000

Scenario A04. Evidence: EV-A04-CC-1, EV-A04-CC-2, EV-A04-CC-3, EV-A04-CC-4 (organisation's controls); EV-A04-ME-1, EV-A04-ME-2, EV-A04-ME-3, EV-A04-ME-4 (independent authority).

This is the difference between **local** safety (is this action acceptable?) and **global** safety (is where these actions lead acceptable?). The same pattern appeared when client money was moved out through an intermediate account, when data was read and then sent out, and when two separately approved role changes combined into a conflict of duties. Controls that check one action at a time cannot see any of these. Each must be tested on its own terms.

What the 48-Hour Audit does

- **Day 0 (scoping call).** Agree one agent or workflow, its tools and permissions, and the outcomes that must never happen.
- **Hours 0-24 (map and attack).** Establish what the agent can actually reach, including tools nobody declared. Then attempt each prohibited outcome deliberately, one action at a time and in sequences.
- **Hours 24-40 (test the controls).** Run each attempt against the organisation's existing controls and against an independent pre-execution authority. Read the results from the systems themselves, not from what the controls report.
- **Hours 40-48 (report).** A board-readable report: what was tested, what was observed, what remains unknown, and what to do next. Every figure is tied to a recorded, replayable decision.

What it does not do

- It does not declare the organisation or the agent safe, and it does not cover agents, tools or channels outside the agreed scope.
- It does not judge whether a particular AI model would misbehave. It tests what happens if it does.
- It is not a production-readiness opinion. That requires a pilot on live systems.

Two engagements

	48-Hour Bounded Exposure Audit	Enterprise Runtime Governance Audit
Proposed fee	$5,000 (proposed, not market-validated)	£40,000-£70,000 (proposed, not market-validated)
Scope	One agent or workflow	Several agents, shared resources and execution boundaries
Answers	Is this agent's exposure controlled, and where not?	Can the organisation evidence control across its agent estate?
Board output	Concise exposure report and next steps (sample Report A)	Assurance report, prioritised remediation, pilot plan (sample Report B)

What can be demonstrated to a prospective customer today

- A live, repeatable run of the demonstration in these pages, in a simulated environment, with every decision replayable and every report figure checked against its evidence.
- Side-by-side results for the same attempts with no control, with typical customer controls, with independent observation only (shadow mode), and with enforcement.
- Honest failure reporting: one case the independent authority did not prevent, one case that depended on configuration, and what fixed them.

Not yet demonstrable: results on a real customer's systems, the operator console (Control Room) view of these findings, and production performance.

ILLUSTRATIVE SAMPLE